

Acknowledgments

If someone contributed to the thesis... might be good to thank them here.

Abstract

An abstract abstracts the thesis...

Kurzzusammenfassung

Hier kommt eine kurze Zusammenfassung der Thesis...
Dies ist ein Mock-Beispiel. Passen Sie den Inhalt an.

Contents

Acknowledgements	i
Abstract	iii
Kurzzusammenfassung	v
1 Introduction	1
1.1 Introduction	1
1.1.1 The Role of Proteomics in Clinical Cancer Research	1
1.1.2 Limitations of Database-Informed Search (DBIS) Strategies	1
1.1.3 De Novo Sequencing: Unbiased Identification of Peptides and PTMs	1
1.1.4 Challenges of Multiplexing in De Novo Sequencing	2
1.1.5 Problem Statement and Research Gap	2
1.1.6 Objectives and Contributions	2
2 Background	5
2.1 Mass Spectrometry-Based Proteomics	5
2.1.1 Bottom-Up Proteomics Workflow	5
2.1.2 Tandem Mass Spectrometry (MS/MS) and Peptide Fragment Ion Theory	5
2.1.3 Tandem Mass Tag (TMT) Labeling	6
2.1.4 Impact of TMT on Fragmentation Patterns	6
2.2 Peptide Identification Strategies	7
2.2.1 Database Search Engines	7
2.2.2 Principles of De Novo Peptide Sequencing	7
2.2.3 Classical De Novo Peptide Sequencing Approaches	8
2.3 Deep Learning and Transformer Models in Proteomics	8
2.3.1 A New Modeling Approach for De Novo Sequencing	8
2.3.2 Transformer Architecture and Self-Attention	9
2.3.3 Spectrum Encoding and Embedding	9
2.3.4 Autoregressive Sequence Generation and Beam Search	9

3 Methods	11
3.1 Fine-tuning Strategy and Model Adaptations	11
3.1.1 Multi-task Learning Architecture	12
3.2 Data Selection and Training Protocol	13
3.2.1 Data Selection and Composition	13
3.2.2 TMT-labeled Dataset	14
3.2.3 Non-TMT Data (Replay Set)	14
3.2.4 Quality Filtering and Pre-processing	15
3.2.5 Training Protocol	16
3.3 Evaluation Strategy and Performance Metrics	16
3.3.1 Confidence Scoring and Peptide Ranking	17
3.3.2 Precision-Coverage Analysis and Stratification	17
3.3.3 Modification-Specific Evaluation	18
3.4 Downstream Validation and Biological Integration	18
3.4.1 Data Acquisition and Preprocessing	18
3.4.2 Peptide Alignment and Sequence Validation	19
3.4.3 Genomic mismatch Evidence (SNPs)	20
3.4.4 Spectral Evidence and Site-Specific Validation	21
4 Results	23
4.1 Model Training and Convergence	23
4.2 Performance Evaluation on the test set	23
4.2.1 Comparative Performance: Non-TMT and Modanovo Baseline	23
4.2.2 Generalization on TMT-Labeled Data	25
4.3 Hyperparameter Optimization and Isotopic Sensitivity	25
4.3.1 Isotope Error Range and Beam Search Synergy	26
4.3.2 Quantitative Impact of Configuration Changes	26
4.3.3 Impact on Specific Modification Classes	26
4.4 Application on independent Glioma TMT Dataset	27
4.4.1 Proteome Alignment and Score Calibration	27
4.4.2 Stratified Performance and Modification Stability	28
4.4.3 Summary of Validation	29
4.5 Uncovering the Dark Proteome: Comparison with MaxQuant	29
4.5.1 Unique Peptide Identifications	30
4.5.2 Identification of Previously Unidentified Spectra	31
4.5.3 From Global Discovery to Specific Variants	31
5 Discussion	33

Supplementary Material	35
5.1 Zusätzliche Abbildungen	35
5.2 Zusätzliche Tabellen	35

x

List of Figures

3.1	Schematic representation of the adapted transformer architecture. The TMT status is fed as a covariate embedding into the decoder alongside the spectral features, as adapted from the Modanovo framework.	13
4.1	Precision-coverage curves across different PTMs. The performance of the final model on TMT-labeled (orange), non-TMT (blue) spectra is shown. Unmodified peptides are shown in light grey for comparison.	24
4.2	Alignment of <i>de novo</i> predictions to the human reference proteome. The proportion of perfect alignments (blue) increases significantly with higher confidence scores. At a score cutoff of > 0.95 , 86.5% of the 398.4k predictions align perfectly.	28
4.3	Proportion of perfectly aligned PSMs stratified by modification type and ranked by confidence score.	29
4.4	Overlap of unique peptide identifications between the proposed <i>de novo</i> model and MaxQuant. The model successfully recovers the majority of MQ identifications while contributing 40,000 additional unique sequences.	30
S1	Histogram of peptide lengths for all data	35
S2	Histogram of peptide lengths for training data (TMT false)	36
S3	Histogram of peptide lengths for validation data (TMT false) . . .	37

List of Tables

4.1 Comparison of model performance across different configurations. 26

1 Introduction

1.1 Introduction

1.1.1 The Role of Proteomics in Clinical Cancer Research

The advent of precision medicine has shifted the focus from broad therapeutic approaches to individualized treatment strategies, particularly in oncology. Proteomics, the large-scale study of proteins and their functions, provides a functional snapshot of the cell that genomics alone cannot capture [Aebersold2016]. In clinical cancer research, specifically in the study of complex malignancies such as gliomas, mass spectrometry (MS)-based proteomics is indispensable for identifying biomarkers and therapeutic targets [Mertins2016]. By analyzing the proteome, researchers can observe the actual effectors of biological processes, making it possible to understand tumor heterogeneity and drug resistance mechanisms at a molecular level.

1.1.2 Limitations of Database-Informed Search (DBIS) Strategies

The standard paradigm for peptide identification is the Database-Informed Search (DBIS). While effective for well-characterized organisms, DBIS faces the “search space challenge.” As the number of considered post-translational modifications (PTMs) and single nucleotide polymorphisms (SNPs) increases, the theoretical search space expands exponentially, leading to a significant loss in statistical power and increased false discovery rates [Chick2015]. Consequently, rare but biologically significant modifications—often crucial in cancer signaling—remain “hidden” because they were not explicitly included in the search parameters.

1.1.3 De Novo Sequencing: Unbiased Identification of Peptides and PTMs

To overcome the constraints of predefined databases, *de novo* peptide sequencing has emerged as a powerful alternative. This approach predicts

the amino acid sequence directly from the MS/MS fragment ions without prior knowledge of the proteome. Recent breakthroughs in deep learning, particularly transformer-based architectures like ModaNovo or Casanovo, have significantly increased the accuracy of these predictions [Yilmaz2022, Rappaport2024]. These models can theoretically identify any peptide sequence, including those with unexpected modifications, making them ideal for discovering novel proteoforms in clinical samples.

1.1.4 Challenges of Multiplexing in De Novo Sequencing

Quantitative proteomics often relies on multiplexing techniques such as Tandem Mass Tags (TMT) to increase throughput and reduce technical variability across clinical cohorts. However, TMT labeling introduces significant complexity into MS/MS spectra. The chemical tags add substantial mass to the N-terminus and lysine residues, and the fragmentation process produces high-intensity reporter ions in the low m/z range [Thompson2003]. These systematic shifts and altered intensity patterns are not well-represented in standard training datasets for *de novo* algorithms, leading to a drop in performance when analyzing multiplexed data.

1.1.5 Problem Statement and Research Gap

Despite the progress in deep learning-based *de novo* sequencing, a significant gap remains: most state-of-the-art models are trained on unlabeled data. When applied to TMT-labeled samples, models like ModaNovo often fail to correctly interpret the mass shifts and the altered fragmentation behavior. There is currently a lack of specialized *de novo* sequencing architectures or fine-tuning strategies that can handle the unique chemical footprint of TMT labeling while simultaneously identifying diverse PTMs. This limitation prevents the full utilization of multiplexed datasets for discovering novel biological insights beyond the reach of database searches.

1.1.6 Objectives and Contributions

The objective of this thesis is to bridge this gap by expanding *de novo* peptide sequencing capabilities to TMT-labeled proteomics data. We propose to evaluate and adapt transformer-based models—specifically focusing on fine-tuning and conditioning strategies—to account for the systematic mass shifts introduced by TMT. By integrating diverse PTMs into this framework, this work aims to provide a robust tool for the unbiased identification of modified peptides in

1.1 Introduction

multiplexed experiments, ultimately uncovering regulatory networks in cancer biology that are overlooked by traditional methods.

1 Introduction

2 Background

In this chapter, the fundamental principles of mass spectrometry-based proteomics and the computational strategies for peptide identification are discussed. Particular focus is placed on the challenges introduced by chemical labeling and the emergence of deep learning models in *de novo* sequencing.

2.1 Mass Spectrometry-Based Proteomics

2.1.1 Bottom-Up Proteomics Workflow

Mass spectrometry (MS)-based proteomics has become the gold standard for the large-scale analysis of proteins in complex biological samples. The most widely adopted strategy is the “bottom-up” approach. In this workflow, proteins are extracted from a biological source and enzymatically digested—typically using trypsin—into smaller peptides before being analyzed by the mass spectrometer [Gevaert2003]. This enzymatic cleavage is essential because peptides are easier to fractionate, ionize, and fragment than intact proteins. Following digestion, the resulting peptide mixture is usually separated by liquid chromatography (LC) and ionized (e.g., via Electrospray Ionization, ESI) to be transferred into the gas phase for analysis [Aebersold2016].

The analysis occurs in two primary stages within the mass spectrometer. In the first stage (MS1), the instrument measures the mass-to-charge ratio (m/z) and intensity of the intact peptides (precursors). This provides a snapshot of the peptide population in a sample at a given time. From this MS1 information, specific precursor ions are selected for the second stage, tandem mass spectrometry (MS2), to extract structural information for identification.

2.1.2 Tandem Mass Spectrometry (MS/MS) and Peptide Fragment Ion Theory

The identification of the amino acid sequence is achieved through Tandem Mass Spectrometry (MS/MS or MS2). In this process, a specific precursor ion is isolated based on its m/z and subsequently subjected to fragmentation [Steen2004]. In high-resolution instruments like the Orbitrap, Higher-energy

2 Background

Collisional Dissociation (HCD) is the preferred method. HCD is a beam-type collision-induced dissociation technique where ions collide with an inert gas (e.g., Nitrogen), leading to internal energy buildup and subsequent bond breakage, producing a predictable pattern of fragment ions.

According to the Roepstorff-Fohlman nomenclature for peptide fragment ion theory, the fragmentation of the peptide backbone occurs primarily at the amide bonds. In the experimental setting of HCD, this gas-phase dissociation predominantly targets the peptide bonds. This results in two main series of ions: b-ions, where the charge remains on the N-terminal fragment, and y-ions, where the charge remains on the C-terminal fragment [Roepstorff2010].

2.1.3 Tandem Mass Tag (TMT) Labeling

Tandem Mass Tag (TMT) labeling is a powerful chemical labeling strategy used for high-throughput multiplexed quantitative proteomics. The TMT molecule is an isobaric tag consisting of three functional groups: a reactive NHS-ester group for covalent attachment to peptide N-termini and lysine side chains, a mass reporter group, and a mass normalizer group [Thompson2003].

Because the tags are isobaric, peptides from different biological samples (up to 18-plex) are labeled, pooled, and appear as a single precursor peak in the MS1 scan [Werner2014]. By enabling the simultaneous analysis of up to 18 samples in a single LC-MS/MS run, TMT labeling minimizes technical batch effects and ensures consistent quantification across channels, significantly reducing the “missing value” challenge compared to label-free workflows [Rauniyar2014]. Furthermore, a “carrier channel”—an isobaric spike-in of a high-abundance proteome—can be used to boost the precursor signal of low-input samples, enhancing identification and sequencing depth in single-cell applications [Budnik2018].

2.1.4 Impact of TMT on Fragmentation Patterns

The use of TMT tags introduces systematic changes to the peptide fragment spectra. Upon fragmentation via HCD, the isobaric tag cleaves at a specific linker region, releasing the low-molecular-weight reporter ions in the m/z 126–135 range for quantification [McAlister2014]. For identification, TMT labeling presents a challenge: the tag adds a constant mass shift to the N-terminus and lysine residues. Furthermore, the presence of the bulky tag can alter the gas-phase basicity and fragmentation efficiency, often leading to different relative intensities of b- and y-ions compared to unlabeled peptides [Hogrebe2018].

As we move from the physical process of generating these spectra, the focus shifts to the computational interpretation of this data, which leads to the different identification approaches.

2.2 Peptide Identification Strategies

2.2.1 Database Search Engines

The most prevalent method for peptide identification is database searching. This strategy relies on a predefined protein sequence database, such as UniProt [TheUniProtConsortium2023]. Computational search engines, including Mascot, SEQUEST, or MaxQuant (Andromeda), perform an *in silico* digestion of these sequences to generate a library of theoretical spectra [Cox2008]. Each experimental MS/MS spectrum is then compared against these theoretical candidates using scoring functions to determine the best match, known as a Peptide-Spectrum Match (PSM) [Eng1994].

While robust, database searching is limited by the predefined “search space.” Although open search strategies exist to identify PTMs without *a priori* definition, traditional closed searches can only identify sequences explicitly included in the database. Consequently, peptides from non-model organisms or those with unexpected biological variants are frequently missed. Including all possible modifications in a search would lead to a combinatorial explosion, drastically increasing false discovery rates (FDR) and computational costs [Nesvizhskii2010].

2.2.2 Principles of De Novo Peptide Sequencing

In contrast to database-driven methods, *de novo* peptide sequencing reconstructs the amino acid sequence directly from the fragment ion peaks in the MS/MS spectrum without any genomic or proteomic reference [Taylor1997]. This approach maps the mass differences between adjacent peaks directly to the masses of amino acids.

By measuring the mass difference between consecutive ions in a series, the corresponding amino acid can be inferred, as each residue (except for the isomers Leucine and Isoleucine) possesses a unique residual mass. For example, a measured mass shift of 113.08 Da identifies Leucine or Isoleucine, whereas a shift of 71.04 Da identifies Alanine [Steen2004]. Historically, however, *de novo* sequencing was limited by spectral noise, low resolution, and incomplete fragmentation.

2.2.3 Classical De Novo Peptide Sequencing Approaches

Early *de novo* peptide sequencing algorithms, such as PEAKS, Novor, and PepNovo, are rooted in rule-based modeling of peptide fragmentation [Ma2003]. These methods typically represent an MS/MS spectrum as a spectrum graph, where nodes correspond to observed m/z values and edges represent mass differences matching specific amino acids. Sequence identification is framed as finding the optimal path through this graph, guided by hand-crafted scoring functions that account for peak intensities and ion types.

However, these classical approaches face significant limitations. They rely heavily on heuristic fragmentation rules which lack flexibility across different mass spectrometry instruments or chemistries. Incomplete fragmentation often leads to “broken” paths in the spectrum graph, while the presence of post-translational modifications (PTMs) or chemical labels like TMT exponentially increases the search space and resulting ambiguity.

2.3 Deep Learning and Transformer Models in Proteomics

The limitations of rule-based *de novo* sequencing approaches motivated the adoption of deep learning methods, which learn peptide fragmentation patterns directly from data. By leveraging large annotated MS/MS datasets, neural networks can model complex, instrument-specific fragmentation behavior and generalize across varying experimental conditions. This data-driven paradigm marked a fundamental shift in *de novo* peptide sequencing.

2.3.1 A New Modeling Approach for De Novo Sequencing

The transition to deep learning enabled the modeling of *de novo* sequencing as a sequence-to-sequence (seq2seq) task, translating spectral peak patterns into amino acid sequences. Convolutional Neural Networks (CNNs) were initially used to extract local spectral features, while Recurrent Neural Networks (RNNs), such as Long Short-Term Memory (LSTM) networks, modeled sequential dependencies between amino acids [Tran2017]. Although these models, like DeepNovo, outperformed classical methods in many settings, they exhibited limitations in capturing long-range dependencies and global spectral context, particularly for longer peptides or spectra with sparse fragmentation.

2.3.2 Transformer Architecture and Self-Attention

The introduction of the Transformer architecture [Vaswani2017] addressed many of these limitations by replacing recurrence with self-attention mechanisms. Self-attention allows the model to assess relationships between all spectral features simultaneously, enabling the integration of complementary evidence such as b- and y-ion pairs distributed across the full m/z range.

In proteomics, this capability is especially beneficial, as peptide evidence is often fragmented and non-local. The Casanova model was a landmark application of Transformers to *de novo* peptide sequencing, employing an encoder-decoder architecture to autoregressively predict peptide sequences [Yilmaz2022]. By combining global spectral context with precursor mass constraints and beam search decoding, Casanova achieved state-of-the-art performance.

2.3.3 Spectrum Encoding and Embedding

A critical component of Transformer-based models is the encoding of MS/MS spectra into suitable input representations. Typically, spectra are transformed into fixed-length embeddings that incorporate m/z values, intensities, and positional information. These embeddings serve as the input tokens for the Transformer encoder, enabling the model to learn fragmentation-aware representations [Qiao2021]. Effective spectrum encoding is essential for robust performance, as it directly influences how well the model can distinguish informative peaks from noise and account for variations in fragmentation patterns.

2.3.4 Autoregressive Sequence Generation and Beam Search

Most modern *de novo* sequencing models generate peptide sequences autoregressively, predicting one amino acid at a time conditioned on previously predicted residues and the encoded spectrum. Beam search is commonly applied during inference to explore multiple high-probability candidate sequences simultaneously.

This strategy allows the model to balance local confidence with global sequence plausibility while enforcing constraints such as precursor mass consistency. As a result, autoregressive decoding with beam search significantly improves identification accuracy, particularly in ambiguous or noisy spectra [Yilmaz2023].

Following Casanova, several models have expanded the capabilities of *de novo* sequencing. DeepNovo pioneered the use of CNNs for feature extraction [Tran2017], while newer architectures like InstaNovo and π -Prime focused on

2 Background

increasing inference speed. To address the limitation of sparse PTM support in early models, Modanovo extended the token vocabulary to include a broad range of amino acid-PTM combinations, demonstrating that Transformers can scale to biologically diverse datasets [**KlaprothAndrade2025**].

Despite these advances, a significant gap remains regarding Tandem Mass Tag (TMT) proteomics.

3 Methods

3.1 Fine-tuning Strategy and Model Adaptations

The approach builds upon the Modanovo framework, a transformer-based architecture designed for the identification of post-translational modifications (PTMs) using experimental spectra [KlaprothAndrade2025].

The transition from unlabeled or label-free spectra to TMT-multiplexed data requires specific adaptations of the underlying deep learning model. In this work, the fine-tuning process involves adjusting the model to recognize TMT labels not as global experimental parameters, but as specific chemical modifications integrated into the sequencing vocabulary.

Tokenization and Vocabulary Expansion

To accommodate TMT labeling, the model's tokenization strategy was expanded. Modanovo utilizes a residue-based vocabulary where each token represents either a standard amino acid or a specific amino acid-PTM combination [KlaprothAndrade2025]. For this study, the configuration was adjusted to include TMT-specific tokens. These tokens account for the fixed mass shifts on N-termini and Lysine (K) residues.

Specifically, the vocabulary was extended by the following residues and their corresponding mass shifts:

- **K[+229.163]**: Lysine with TMT10/16 label.
- **[+229.163]-**: TMT10/16 label at the peptide N-terminus.
- **K[+343.206]**: Lysine with both TMT and GlyGly (ubiquitination) modification.
- **K[+271.173]**: Lysine with both TMT and Acetyl modification.
- **K[+243.179]**: Lysine with both TMT and Methyl modification.

Following the Modanovo initialization protocol, the embeddings for these new tokens were initialized by averaging the embeddings of their

3 Methods

constituent components (e.g., the base amino acid embedding and the modification-specific shift) to leverage pre-learned chemical representations [KlaprothAndrade2025].

TMT Covariate Embedding

To enable the decoder to account for systematic shifts in fragmentation patterns and physicochemical properties induced by TMT labeling, we introduce a categorical conditioning mechanism. This allows the model to explicitly distinguish between TMT-labeled and unlabeled spectra at a global level.

Analogous to the embedding of precursor features (precursor mass and charge), we define a learnable TMT-specific embedding. For each spectrum, a binary indicator $f_{\text{TMT}} \in \{0, 1\}$ encodes the presence or absence of TMT labeling and is mapped through an embedding layer:

$$\mathbf{E}_{\text{TMT}} = \text{Embedding}(f_{\text{TMT}}) \in \mathbb{R}^{d_{\text{model}}}.$$

The resulting vector is integrated into the latent representation via additive fusion. Specifically, it is added to the precursor embedding prior to decoding:

$$\text{prec_emb}_{\text{conditioned}} = \text{prec_emb} + \mathbf{E}_{\text{TMT}}.$$

By injecting this information at the level of the precursor representation—effectively seeding the start of the decoding process—the transformer can adapt its internal representations to the chemical environment associated with TMT-labeled peptides. This conditioning strategy is computationally efficient, as it preserves the model dimensionality while providing a strong global signal that guides de novo sequencing depending on the labeling state of the sample.

3.1.1 Multi-task Learning Architecture

Building on the modularity of Modanovo, a multi-task learning head was evaluated. This architectural extension aims to decouple the prediction of the amino acid backbone from the specific PTM state by utilizing a dedicated PTM prediction head [KlaprothAndrade2025].

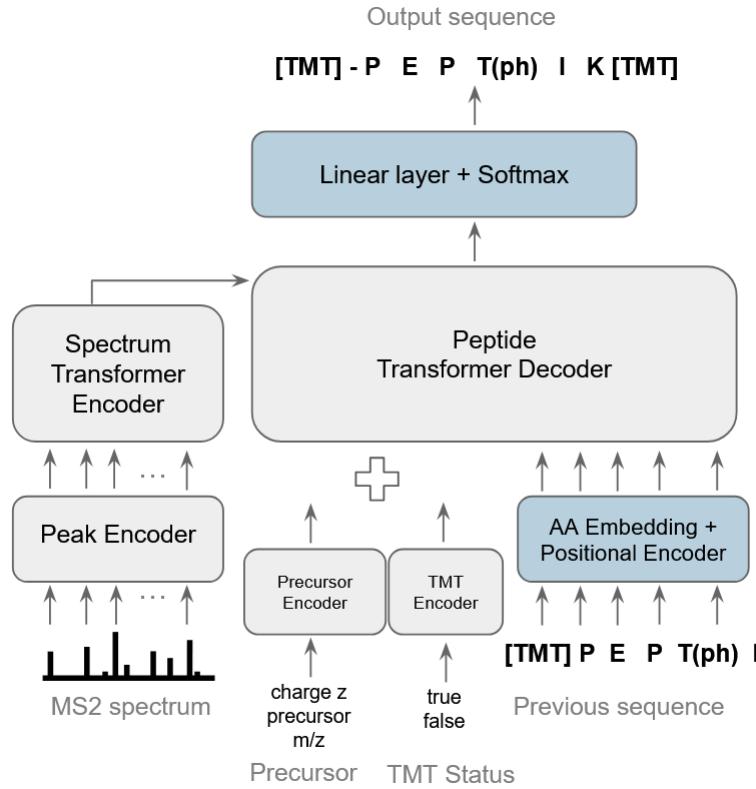


Figure 3.1: Schematic representation of the adapted transformer architecture. The TMT status is fed as a covariate embedding into the decoder alongside the spectral features, as adapted from the Modanovo framework.

3.2 Data Selection and Training Protocol

3.2.1 Data Selection and Composition

For the training and evaluation of the adapted model, a robust dataset was curated to ensure high-quality spectral representations. The fundamental requirement for supervised learning in this context is the availability of ground truth sequences associated with high-resolution fragment spectra. Data were integrated from various sources, initially stored in CSV and mzML formats, and subsequently compiled into a unified Mascot Generic Format (MGF) file. This format allows for a streamlined input pipeline where the peptide sequence is explicitly linked to its corresponding spectrum [Deutsch2012].

3.2.2 TMT-labeled Dataset

The primary dataset for adapting the model to isobaric labeling is derived from the PROSPECT-MultiPTM collection [Zeng2024]. These data are based on the ProteomeTools project, a large-scale synthetic peptide library effort [Zolg2017].

Instrumentation and Fragmentation All spectra were acquired using Thermo Scientific Orbitrap instruments (Q Exactive and Orbitrap Fusion series). Fragmentation was performed exclusively using Higher-energy Collisional Dissociation (HCD), resulting in high-resolution MS₂ spectra. The raw data are hosted via the PRIDE archive and ProteomeXchange [Perez-Riverol2022].

Labeling and Modifications The peptides in this dataset are labeled with Tandem Mass Tags (TMT). The chemical modification manifests as a specific mass shift at the peptide N-terminus and on the ϵ -amino group of all Lysine (Lys) residues. The exact mass shifts follow the Unimod definitions and are encoded using the ProForma standard [Leis2022].

Dataset Statistics The TMT-labeled dataset is partitioned into training, validation, and test sets with the following dimensions:

- **Training Set:** 3,683,888 Peptide-Spectrum Matches (PSMs) covering 75,268 unique peptides.
- **Validation Set:** 363,612 PSMs covering 7,751 unique peptides.
- **Test Set:** 364,867 PSMs covering 9,415 unique peptides.

3.2.3 Non-TMT Data (Replay Set)

To prevent catastrophic forgetting during the fine-tuning process, a diverse reference dataset of unlabeled (non-TMT) spectra is included. This “Replay Set” consists of a mixture of 80% MultiPTM data and 20% data from the MassIVE Knowledge Base (MassIVE-KB) [Wang2018].

Technical Characteristics The instrumentation and fragmentation settings (Orbitrap HCD) are consistent with the TMT-labeled dataset to ensure technical compatibility. This set includes a wide variety of post-translational modifications and biological sequences.

Dataset Statistics The non-TMT reference data is partitioned into training, validation, and test sets with the following dimensions:

- **Training Set:** 784,128 PSMs covering 289,568 unique peptides.
- **Validation Set:** 98,396 PSMs covering 23,004 unique peptides.
- **Test Set:** 93,453 PSMs covering 19,141 unique peptides.

3.2.4 Quality Filtering and Pre-processing

To minimize noise and prevent the model from learning experimental artifacts, several quality filtering steps were applied:

- **Peak Cleaning:** Spectra containing no intensity information or empty peaks were removed to maintain data density.
- **Distribution Analysis:** The dataset was screened for biases in peptide length and charge state distributions. Ensuring a representative spread across these parameters is crucial for the generalization of transformer-based models [KlaprothAndrade2025].
- **Bias Mitigation:** To prevent the model from overfitting to hyper-abundant peptides, a threshold of 229 Peptide-Spectrum Matches (PSMs) per unique peptide sequence was enforced.
- **Data Leakage Prevention:** Following the rigorous validation protocols of *ModaNovo*, a strict data split was implemented. Peptides with specific modifications (e.g., *PEP[ph]*) were assigned to the same split as their unmodified counterparts (*PEP*) to ensure that the model learns the chemical principles of modifications rather than memorizing specific sequences [KlaprothAndrade2025].

The final dataset was structured into an 80/10/10 split (training, validation, and testing). To specifically address the TMT expansion, an 80/20 balance between TMT-labeled and unlabeled spectra was maintained, and all non-TMT spectra originating from TMT-specific experiments were removed to ensure label consistency. Furthermore, Unimod syntax was translated into mass-shift syntax (e.g., [+229.163]) to align with the model's vocabulary.

3.2.5 Training Protocol

Model fine-tuning was initialized from the publicly available ModaNovo checkpoint, allowing the model to build upon previously learned representations of peptide fragmentation and spectral structure [KlaprothAndrade2025]. In contrast to partial adaptation strategies, all model parameters were updated during fine-tuning, i.e. no layers were frozen, enabling global adaptation to TMT-specific fragmentation effects and modification patterns.

The underlying transformer architecture was kept identical to the base ModaNovo configuration, comprising a model dimension of $d_{\text{model}} = 512$, 8 self-attention heads, a feed-forward dimension of 1024, and 9 layers each in the encoder and decoder stacks. This architectural consistency ensures that any observed performance differences can be attributed to the fine-tuning procedure rather than structural changes.

Optimization hyperparameters were deliberately chosen to favor stable adaptation of the pre-trained weights. A low learning rate of 1×10^{-6} was used to prevent catastrophic forgetting while still permitting gradual adjustment to TMT-induced shifts in fragmentation behavior. To further stabilize early training dynamics, a warm-up phase of two epochs was applied. Regularization was introduced via a weight decay of 1×10^{-5} and label smoothing with a factor of 0.01, improving generalization in the presence of heterogeneous modification patterns.

Training was performed using mixed-precision arithmetic with *bf16* precision, reducing memory consumption and improving computational efficiency on modern GPU architectures without compromising numerical stability [Micikevicius2017]. Model selection was based on validation loss, and the checkpoint with the lowest validation loss was retained for all downstream analyses.

3.3 Evaluation Strategy and Performance Metrics

To rigorously assess the performance of the TMT-adapted de novo sequencing model, a multi-faceted evaluation framework was established. The primary objective is to determine how well the model generalizes to TMT-labeled spectra and various post-translational modifications (PTMs) compared to traditional database-driven assignments.

3.3.1 Confidence Scoring and Peptide Ranking

Each peptide-spectrum match (PSM) generated by the model is assigned a confidence score to facilitate ranking and quality control. Following the architecture of Transformer-based models like Casanova and ModaNovo, we derive a peptide-level score by calculating the arithmetic mean of the individual amino acid confidence scores, which are obtained from the softmax output at each decoding step [Yilmaz2022].

To ensure the physical plausibility of the predictions, a mass-matching constraint is applied. If the calculated mass of the predicted sequence (including PTMs and TMT labels) deviates from the observed precursor mass beyond a defined tolerance (e.g., 10 ppm), the peptide score is penalized. This integration of spectral evidence and thermodynamic constraints is crucial for distinguishing between high-confidence sequences and plausible but incorrect mass-shift combinations.

3.3.2 Precision-Coverage Analysis and Stratification

The core metric for evaluating the model’s predictive power is the precision-coverage curve. This allows for a threshold-independent assessment of how many peptides can be identified at a given reliability level. For this study, we specifically focus on the Area Under the Precision-Coverage Curve (AUPCC), calculated using the trapezoidal rule [Pedregosa2011].

A critical aspect of our evaluation is the **stratified analysis**. To understand the specific impact of the TMT expansion, we evaluate the performance separately for:

- **TMT-labeled spectra:** To measure the success of the model adaptation to systematic mass shifts.
- **Unlabeled (non-TMT) spectra:** To ensure that the model retains its general sequencing capabilities without losing performance on standard data (preventing “catastrophic forgetting”).

Precision (P) and Coverage (C) at a score threshold t are defined as:

$$P(t) = \frac{|\text{Correct PSMs with score } \geq t|}{|\text{Total predictions with score } \geq t|} \quad (3.1)$$

$$C(t) = \frac{|\text{Predictions with score } \geq t|}{|\text{Total ground truth identifications}|} \quad (3.2)$$

3 Methods

A PSM is considered correct if the sequence exactly matches the ground truth identified by a database search (e.g., MaxQuant or MSFragger), treating isobaric amino acids such as Leucine and Isoleucine as equivalent [KlaprothAndrade2025].

3.3.3 Modification-Specific Evaluation

To uncover biological insights beyond standard searches, we evaluate the precision for specific PTM-amino acid combinations. For a given modification (e.g., Phosphorylation at T), we subset the ground truth data to include all peptides containing this specific shift. This granular view ensures that the model's ability to handle complex, multiplexed PTM patterns is validated across both TMT and non-TMT backgrounds.

3.4 Downstream Validation and Biological Integration

To evaluate the practical utility of the model, it was applied to an independent Glioma cancer dataset.

3.4.1 Data Acquisition and Preprocessing

The primary data source consisted of Orbitrap-based mass spectrometry raw files (.raw) from cancer proteomic studies. To ensure compatibility with the deep learning framework, raw files were converted into .mgf format using the ThermoRawFileParser (v2.0.0), which facilitates the extraction of metadata and spectral information [Hulstaert2020]. All MS1 (precursor scans) and MS3 scans were removed for the analysis.

De Novo Sequencing Configuration (finetuned ModaNovo) The sequencing of the TMT-labeled spectra was performed using the adapted ModaNovo framework. To ensure high-quality sequence predictions and to accommodate the systematic shifts introduced by TMT, the following parameters were applied:

- **Mass Tolerances:** The precursor mass tolerance was set to 50 ppm to account for potential drift in large-scale datasets. The isotope error range was restricted to [0, 3].

- **Spectrum Processing:** To reduce noise, only the 150 most intense peaks per spectrum (*n_peaks*) within an *m/z* range of 50.0 to 2500.0 were retained. Peaks within a 2.0 Da window of the precursor *m/z* were removed to prevent interference.
- **Sequence Constraints:** A minimum peptide length of 6 amino acids and a maximum of 100 were enforced. For the decoding process, a beam search with a width of *n_beams* = 1 was utilized, focusing on the top-ranked match.
- **Quality Control:** Only spectra with a precursor charge of ≤ 10 and a relative intensity threshold of 0.01 were processed.

Database Search Configuration (MaxQuant) To provide a complementary ground-truth baseline, all Glioma datasets were processed using MaxQuant (version 2.1.3.0) [Tyanova2016]. The search was conducted against the human reference proteome (UniProt UP000005640).

The search parameters were harmonized with the experimental design:

- **Protease:** Trypsin/P was specified, allowing for cleavage C-terminal to Lysine and Arginine, even when followed by Proline.
- **Fixed Modifications:** Carbamidomethylation of cysteine (+57.021 Da) was set as a static modification.
- **Variable Modifications:** To capture the regulatory landscape of the cancer samples, oxidation (M) and phosphorylation (S/T/Y) were included in the search space.

The results from this database-driven approach serve as the benchmark for calculating the precision and coverage of the de novo model's predictions.

3.4.2 Peptide Alignment and Sequence Validation

To validate the biological origin of the predicted de novo sequences and to distinguish between known peptides and potential novel discoveries, a sequence alignment against the reference proteome was performed.

BLASTp Configuration and Ambiguity Handling Alignment was executed using Protein-Protein BLAST (version 2.17.0+) against the same human reference proteome (UniProt UP000005640) used for the MaxQuant search. To account for the inherent mass-spectrometric ambiguities where different amino

3 Methods

acid compositions result in near-identical masses, a specialized encoding strategy was implemented:

- **I/L Equivalence:** Leucine and Isoleucine were treated as identical.
- **Mass-Based Ambiguities:** Specific residues with overlapping mass shifts, such as the deamidation of Glutamine ($Q + 0.984$ Da) resulting in the same mass as Glutamic acid (E), or Asparagine deamidation matching Aspartic acid (D), were addressed.
- **Ambiguity Codes:** All unmodified Aspartic acid (D) residues in the query were replaced with the code B (representing D or N), and all unmodified Glutamic acid (E) residues were replaced with Z (representing E or Q). The PAM30 substitution matrix was utilized, as it natively handles these ambiguity codes to allow perfect matches against either potential target residue.

The alignment parameters were set to an E-value of 2000, qcov_hsp_perc of 80, and comp_based_stats disabled to prevent score inflation for short peptide sequences.

Post-Alignment Filtering and SNP Analysis The resulting alignments were further enriched to identify Single Nucleotide Polymorphisms (SNPs) and truncation events. To account for sequences extending beyond protein termini or alignment gaps, the full query and target sequences were retrieved. The number of mismatches was calculated by considering the B and Z equivalences. For each remaining mismatch, an automated codon-lookup was performed. A mismatch was classified as “SNP-explainable” if the transition between the predicted amino acid and the reference residue could be achieved by a single nucleotide substitution in the underlying codon. Cases involving gaps or truncations where the query extended beyond the reference boundaries were excluded from the SNP analysis to maintain high confidence in the mutation mapping.

3.4.3 Genomic mismatch Evidence (SNPs)

To provide orthogonal validation for the identified peptide variants, de novo predictions were cross-referenced with patient-specific genomic data. This integration aimed to confirm whether non-canonical peptides identified by the model correspond to somatic mutations detected via DNA sequencing.

Genomic Variant Processing Genomic data from patient-matched tumor samples were obtained through panel sequencing of approximately 500 cancer-relevant genes. The raw variants were annotated using the Nirvana pipeline (version 3.2.3) based on the GRCh37 (hg19) genome assembly. Functional annotation was performed via the Variant Effect Predictor (VEP), integrating data from ClinVar, COSMIC, dbSNP, and gnomAD. For this study, the analysis focused on somatic mutations with a MODERATE or HIGH functional impact, specifically prioritizing missense variants and in-frame insertions/deletions (indels).

Mapping and Cross-Validation Pipeline A systematic mapping procedure was implemented to link mass spectrometric evidence to genomic coordinates:

- **ID Mapping:** UniProt identifiers obtained from the BLASTp alignments (see Section 4.4.1) were mapped to Ensembl transcript IDs used in the VEP-annotated genomic files to ensure data consistency.
- **Positional Verification:** For every de novo peptide that aligned to a cancer-related gene with SNP-explainable mismatches, the exact amino acid position within the protein sequence was determined.
- **Evidence Integration:** We cross-checked whether the specific amino acid substitution predicted by the de novo model (and categorized as SNP-explainable by the codon-lookup) matched the genomic evidence (HGVSp/HGVSp notations) provided in the patient’s Annotated.json files.

A peptide was considered “genomically validated” only if the de novo sequence matched the specific mutation site and the substituted amino acid identified in the patient’s mutational landscape, thereby bridging the gap between mass spectrometry-based discovery and clinical genomics.

3.4.4 Spectral Evidence and Site-Specific Validation

To ensure the reliability of identified Single Amino Acid Variants (SAAVs) and post-translational modifications, a stringent spectral validation framework was implemented. This went beyond simple sequence matching by analyzing the fragmentation patterns and the recurrence of specific hits across the dataset.

Validation of Single Nucleotide Polymorphisms (SNPs) For every peptide sequence where a mismatch was explained by a genomic SNP, the following metrics were evaluated to confirm the discovery:

3 Methods

- **PSM Abundance and Patient Coverage:** We quantified the total number of Peptide-Spectrum Matches (PSMs) supporting the specific SNP and assessed its prevalence across the patient cohort.
- **Diagnostic Ion Intensity:** The presence and relative intensity of “site-determining” fragment ions were analyzed. Specifically, we verified that the *b*- and *y*-ion series flanking the mutation site (e.g., y_n and y_{n-1} for a C-terminal mutation) provided unambiguous evidence for the substituted amino acid.
- **Comparative Identification:** For each scan supporting a SNP, the corresponding MaxQuant (MQ) assignment was retrieved. This allowed us to distinguish between cases where MQ had a “near-match” or where the spectrum remained entirely unassigned in the database-driven search.

Phosphorylation Site (P-Site) Verification For peptides carrying phosphorylation (S/T/Y), we localized the modification to specific residues and calculated their absolute position within the protein sequence. The quality of these assignments was assessed through:

- **Site Occupancy and Redundancy:** We determined how many PSMs identified the modified peptide and compared this to the number of PSMs where the same peptide sequence was found in its unmodified state.
- **Sequence Redundancy:** To increase confidence, we investigated if the same phosphorylation site was validated by different peptide species (e.g., peptides resulting from missed cleavages or different precursor charge states).
- **Fragmentation Pattern Analysis:** Similar to SNP validation, we focused on the intensity of fragment ions carrying the +79.966 Da mass shift. The presence of characteristic neutral losses (e.g., -98 Da for Phospho-S/T) and the specific *b/y* ion coverage around the P-site were used to confirm the site localization.

This multi-layered approach ensures that the identified PTMs and SAAVs are not artifacts of the de novo model but are supported by high-quality experimental evidence.

4 Results

4.1 Model Training and Convergence

The model was trained until the validation loss showed early signs of stagnation, indicating the onset of overfitting. Convergence was reached after 13 epochs, corresponding to approximately 1,550,000 training steps. To assess the final performance, the model state at this checkpoint was utilized for inference on the complete, independent test dataset.

4.2 Performance Evaluation on the test set

To evaluate the predictive confidence, precision-coverage curves were generated for the different PTM categories and stratified by TMT vs non-TMT data. The performance is quantified using the Area Under the Precision-Coverage Curve (AUPCC).

The evaluation follows a peptide-centric approach: a single ground truth peptide contributes to the curves of all modifications it contains. For instance, a peptide sequence such as “[+229.997]-PEPT[+79.966]IDEK[+14.016]” is included in both the precision-coverage curve for phosphorylated threonine (T[+79.966]) and monomethylated lysine (K[+14.016]). In all subsequent analyses, the performance on unmodified peptides (light grey in Figure 4.1) serves as a reference baseline.

4.2.1 Comparative Performance: Non-TMT and Modanovo Baseline

A key requirement for the expanded model was to maintain high performance on standard (non-TMT) data, ensuring that the addition of TMT-related features did not degrade general sequencing accuracy. On non-TMT unmodified peptides, the model achieved an AUPCC of 0.92, which is highly consistent with the performance reported for the Modanovo architecture (AUPCC of 0.93) [Yilmaz2022].

4 Results

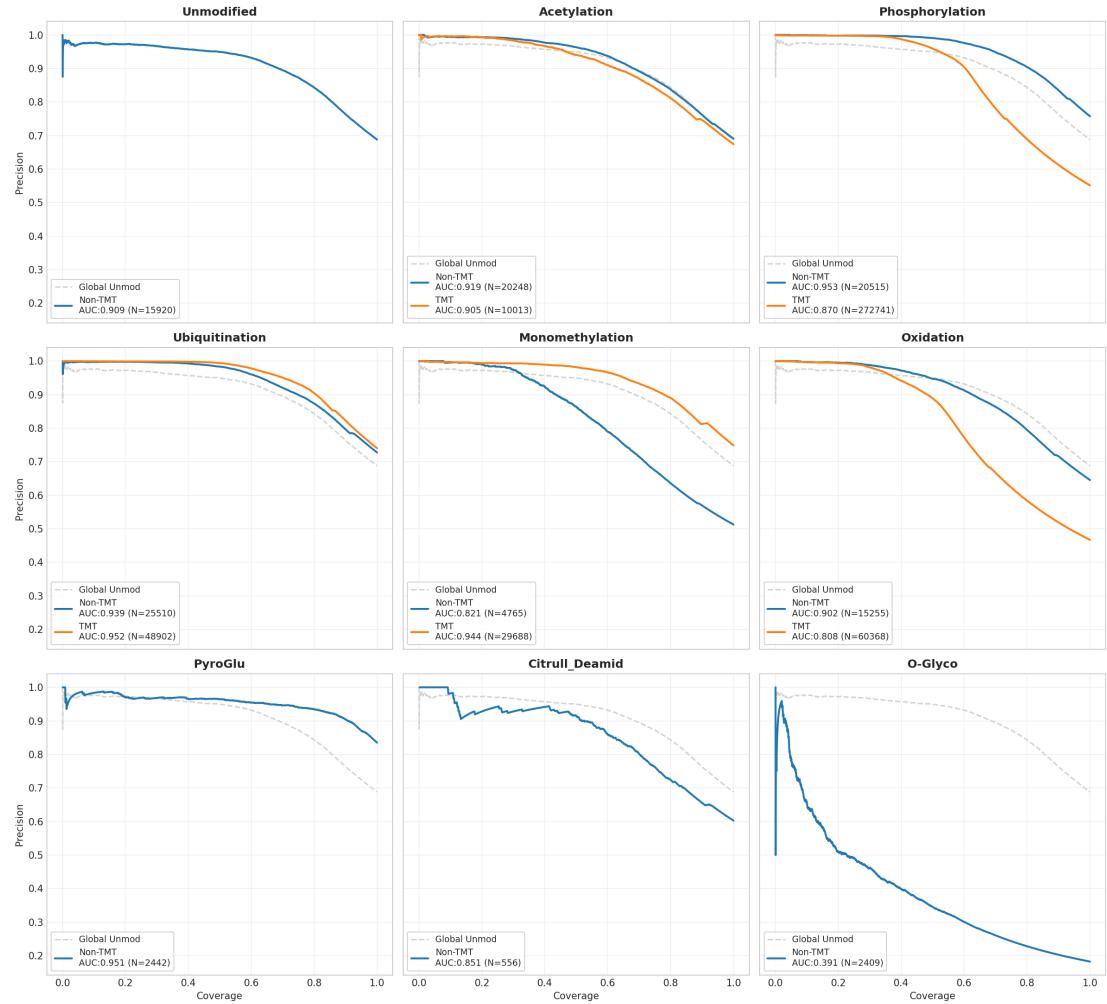


Figure 4.1: Precision-coverage curves across different PTMs. The performance of the final model on TMT-labeled (orange), non-TMT (blue) spectra is shown. Unmodified peptides are shown in light grey for comparison.

This stability is further reflected in the specific PTM performance. For instance, the model demonstrated similar accuracy for Phosphorylation on non-TMT data (AUPCC 0.95) staying aligned with the general Modanovo median for phosphorylation (0.93–0.96 depending on the residue) [Yilmaz2022]. Similarly, high performance was observed for Pyro-Glu (AUPCC 0.95), likely due to its restricted occurrence at the peptide N-terminus, a pattern the model successfully internalizes. In contrast, O-Glycosylation (S/T[+203.079]) remains a significant challenge with an AUPCC of 0.39, mirroring the difficulties reported in literature due to low training examples and complex fragmentation [Yilmaz2022].

4.2.2 Generalization on TMT-Labeled Data

The global performance on TMT data (AUPCC 0.89) is slightly lower than on non-TMT data (AUPCC 0.92), reflecting the increased spectral complexity introduced by the chemical labels.

Specific findings for TMT-labeled PTMs include:

- **Ubiquitination and Monomethylation:** Interestingly, the model performed slightly better on TMT-labeled Ubiquitination (AUPCC 0.95) compared to its non-TMT counterpart (AUPCC 0.94). For Monomethylation, the gap was even more pronounced (TMT: 0.94 vs. non-TMT: 0.82), suggesting that the systematic mass shifts provided by TMT might aid in distinguishing these specific modifications in certain contexts.
- **Acetylation:** Performance remained robust across both categories (TMT: 0.91 vs. non-TMT: 0.92), indicating that Acetylation is identified with high confidence regardless of the labeling strategy.
- **Oxidation:** A noticeable decrease in performance was observed for TMT-labeled oxidized peptides (AUPCC 0.81) compared to non-TMT (AUPCC 0.90), pointing towards potential interference between the TMT-label fragments and the neutral losses associated with methionine oxidation.

The final precision at full coverage (e.g., 0.59 on all tmt-labeled spectra) demonstrates that the model provides a reliable foundation for uncovering biological insights in multiplexed quantitative proteomics experiments.

4.3 Hyperparameter Optimization and Isotopic Sensitivity

A critical challenge identified during the evaluation of TMT-labeled data was the increased frequency of non-monoisotopic precursor selection. In TMT mul-

4 Results

triplexed samples, the high density of ions and the chemical labeling itself often lead to a shift where the instrument triggers fragmentation on an isotopic peak rather than the monoisotopic mass. To address this, the model configuration was iteratively optimized.

4.3.1 Isotope Error Range and Beam Search Synergy

The initial baseline configuration (Config0) utilized a restrictive isotope error range of 0, 1, which often forced the model to strictly adhere to the theoretical monoisotopic mass, leading to sequence errors when the precursor mass was incorrectly assigned by the instrument. By expanding this range to 0, 3 (Config ISO), the model gained the necessary flexibility to account for these systematic mass shifts without being penalized for “off-by-one” or “off-by-two” Dalton errors.

Furthermore, while literature for models like Casanova suggests that increasing the beam size has diminishing returns, our results indicate a strong synergy between isotope flexibility and search breadth. Increasing the beam size from 1 to 5 (Config FINAL) allowed the decoder to explore a wider sequence space, which proved essential for resolving the complex fragmentation patterns of TMT-labeled and modified peptides.

4.3.2 Quantitative Impact of Configuration Changes

The optimization process led to a stepwise improvement in all global metrics (see Table 4.1). The global AUPCC increased from 0.8746 (Config0) to 0.8817 (ISO) and reached its peak at 0.8988 in the FINAL configuration. Notably, the most significant breakthrough occurred in the jump to the FINAL configuration, where the global precision at full coverage rose from 0.5726 to 0.6146.

Table 4.1: Comparison of model performance across different configurations.

Metric	Config0	Config ISO	Config FINAL
Global AUPCC	0.8746	0.8817	0.8988
Global Precision (Cov 1.0)	0.5726	0.5726	0.6146

4.3.3 Impact on Specific Modification Classes

The transition to Config FINAL particularly benefited complex PTMs:

- **Ubiquitination:** This category showed a massive precision jump from 0.6675 to 0.7359. The combination of expanded isotope ranges and increased beam size likely allowed the model to better identify the characteristic GlyGly-remnant fragmentation patterns, which are often obscured in TMT-labeled spectra.
- **Phosphorylation:** As the largest dataset (N=293,256), the shift from 0.8556 (Config0) to 0.8784 (AUPCC) and a precision increase to 0.5656 is statistically the most significant indicator of the model’s enhanced robustness.
- **Monomethylation:** The AUPCC rose from 0.9032 to 0.9326. Remarkably, TMT-labeled methylated peptides achieved a significantly higher AUPCC (0.9443) than their non-TMT counterparts (0.8211). This suggests that the model learned to leverage the TMT-induced mass shifts as a “fingerprint” to distinguish methylation from isobaric interferences more effectively than in unlabeled data.

The improvement in Global TMT AUPCC (from 0.8640 to 0.8717 in the ISO stage) confirms that accounting for isotopic uncertainty is a prerequisite for high-confidence *de novo* sequencing in multiplexed workflows.

4.4 Application on independent Glioma TMT Dataset

To evaluate the model’s discovery potential, inference was performed on the complete Glioma TMT dataset, comprising 6.48 million spectra. Given that *de novo* sequencing operates independently of protein databases, it is crucial to validate the predicted sequences against a reference proteome to distinguish between high-confidence identifications and potential false positives.

4.4.1 Proteome Alignment and Score Calibration

The predicted sequences for unmodified peptides were aligned against the human reference proteome using `blastp`. This alignment process serves as a first-tier validation of the model’s output quality. The sequences were categorized based on their alignment criteria: perfect matches (identity = 100%), sequences with one mismatch, and sequences with two mismatches.

Out of the total predictions, 1.34 million peptides aligned perfectly with the reference proteome, while 0.62 million and 0.8 million sequences exhibited one and two mismatches, respectively.

4 Results

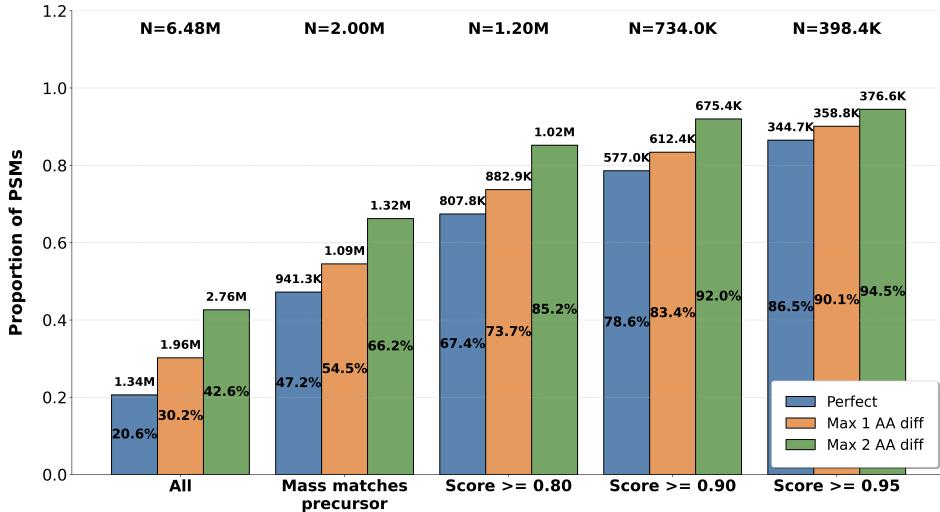


Figure 4.2: Alignment of *de novo* predictions to the human reference proteome. The proportion of perfect alignments (blue) increases significantly with higher confidence scores. At a score cutoff of > 0.95 , 86.5% of the 398.4k predictions align perfectly.

The upward trend of the blue bars in Figure 4.2 confirms that the internal scoring mechanism of the Transformer architecture effectively reflects the probability of a sequence being biologically “correct.” At the highest confidence bin (> 0.95), the vast majority of sequences are already known to the genome, providing a solid baseline for the subsequent analysis of modified and novel peptides.

4.4.2 Stratified Performance and Modification Stability

To further investigate the reliability across different chemical states, the precision of the predictions was stratified by modification type and ranked by confidence (PSM Rank). In this analysis, the “Rank 1” prediction represent the highest-confidence spectrum.

As shown in Figure 4.3, most modifications follow the expected trend where decreasing confidence ranks correlate with a lower proportion of perfect alignments. Phosphorylated peptides follow the global trend closely.

A notable outlier is Monomethylation. While this modification showed excellent performance in the AUPCC metrics on the test set, the proportion of perfect alignments drops more sharply with increasing rank compared to other modifications. This discrepancy suggests the absolute mass shifts or the localization

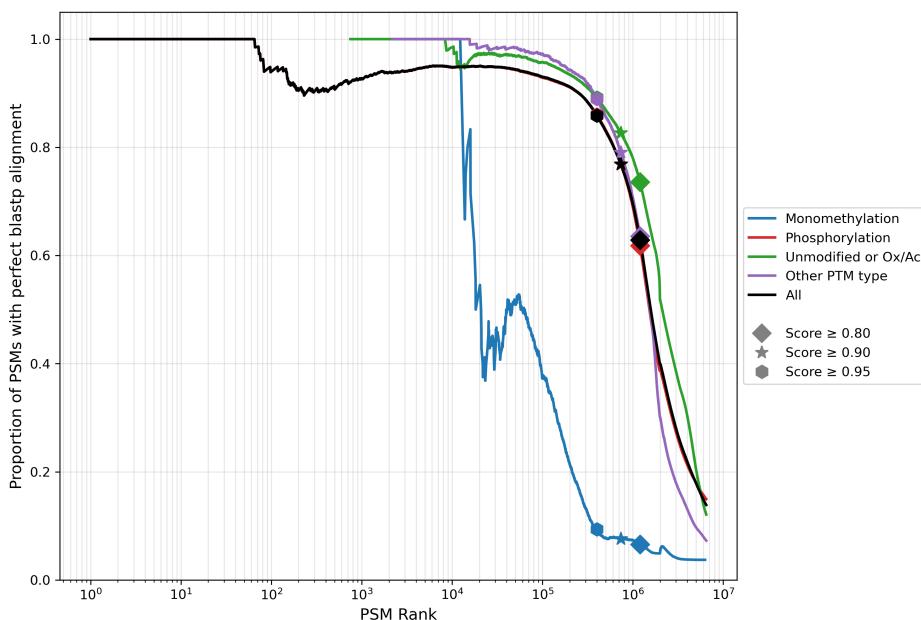


Figure 4.3: Proportion of perfectly aligned PSMs stratified by modification type and ranked by confidence score.

of the methyl group might still lead to mismatches during genomic alignment in real-world samples.

4.4.3 Summary of Validation

The alignment results demonstrate that the expanded model produces biologically plausible peptide sequences. The high correlation between the model's confidence score and the genomic match rate validates the use of these scores as a filter for discovery. This foundation allows for the exploration of spectra that remain unidentified by traditional database-driven methods, such as MaxQuant, which will be discussed in the following section.

4.5 Uncovering the Dark Proteome: Comparison with MaxQuant

A primary objective of this study was to evaluate the extent to which *de novo* peptide sequencing can expand the identification landscape beyond the limitations of database-driven methods. By comparing our results with those obtained via MaxQuant (MQ), we can quantify the “dark matter” of the pro-

teome—spectra that contain high-quality information but remain unidentified by traditional search engines.

4.5.1 Unique Peptide Identifications

The comparison of unique peptide sequences identified by both methods reveals a substantial expansion of the detectable proteome. As shown in Figure 4.4, the *de novo* approach identified approximately 40,000 unique peptides that were also found by MaxQuant, demonstrating high consistency in the “known” space.

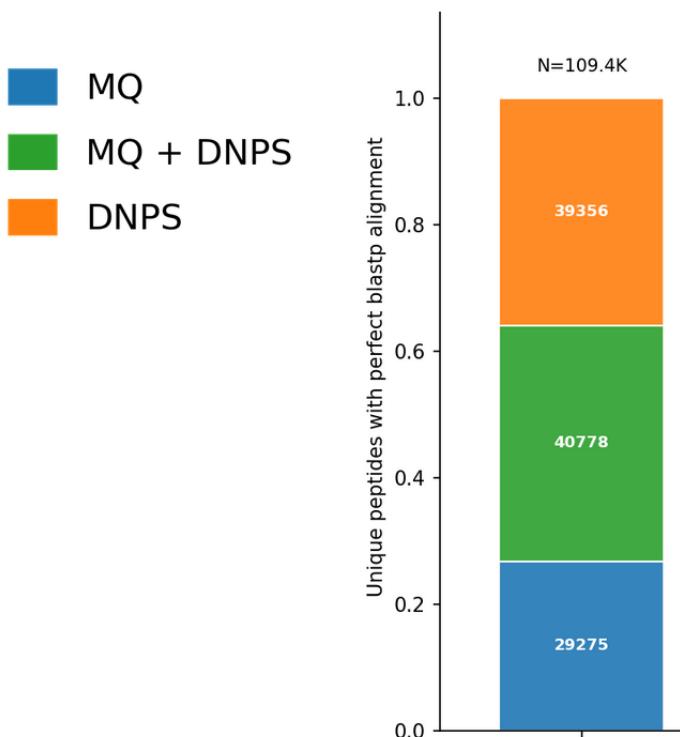


Figure 4.4: Overlap of unique peptide identifications between the proposed *de novo* model and MaxQuant. The model successfully recovers the majority of MQ identifications while contributing 40,000 additional unique sequences.

Crucially, the model discovered an additional 40,000 unique peptides that were entirely absent from the MQ results. Conversely, only a negligible fraction (approximately 30 peptides) was identified by MQ but missed by our model, indicating that the *de novo* approach covers nearly the entire search space of the database-driven method while doubling the number of unique sequences.

4.5.2 Identification of Previously Unidentified Spectra

The true potential of the model is reflected at the spectrum level (PSMs). In the analyzed dataset, MaxQuant provided identifications for 1,510,866 spectra. Our model was able to provide high-confidence sequences for a significant portion of the remaining “dark” spectra:

- **High-Confidence Matches (Score > 0.8):** We identified 446,290 additional spectra with a confidence score above 0.8. Based on our previous alignment validation, this score range corresponds to highly reliable peptide sequences.
- **Precursor-Shifted Matches (Score $-0.2 < s < 0$):** Interestingly, we found 231,208 spectra in a score range that indicates high sequencing confidence but carries a penalty (-1) due to a precursor mass mismatch. These spectra likely represent peptides with unexpected modifications or amino acid substitutions (SNPs) that shift the precursor mass beyond the tolerance of the theoretical database entry, yet yield clear fragmentation patterns.

Summing these categories, the *de novo* approach provides plausible sequence candidates for over 670,000 spectra that were previously discarded in the MQ workflow. This substantial increase in spectral utilization highlights the model’s ability to move beyond the “closed-search” paradigm.

4.5.3 From Global Discovery to Specific Variants

The identification of these additional sequences suggests that the “dark proteome” in these glioma samples is rich in biological variants. The high number of high-confidence predictions that do not perfectly match the database entries—especially those with slight precursor shifts—points towards the presence of non-canonical protein isoforms. In the following sections, we will categorize these findings into specific biological phenomena, namely novel phosphorylation sites (p-sites) and single nucleotide polymorphisms (SNPs).

4 Results

5 Discussion

5 Discussion

Supplementary Material

Dieses Kapitel enthält zusätzliches Material zur Thesis.

5.1 Zusätzliche Abbildungen

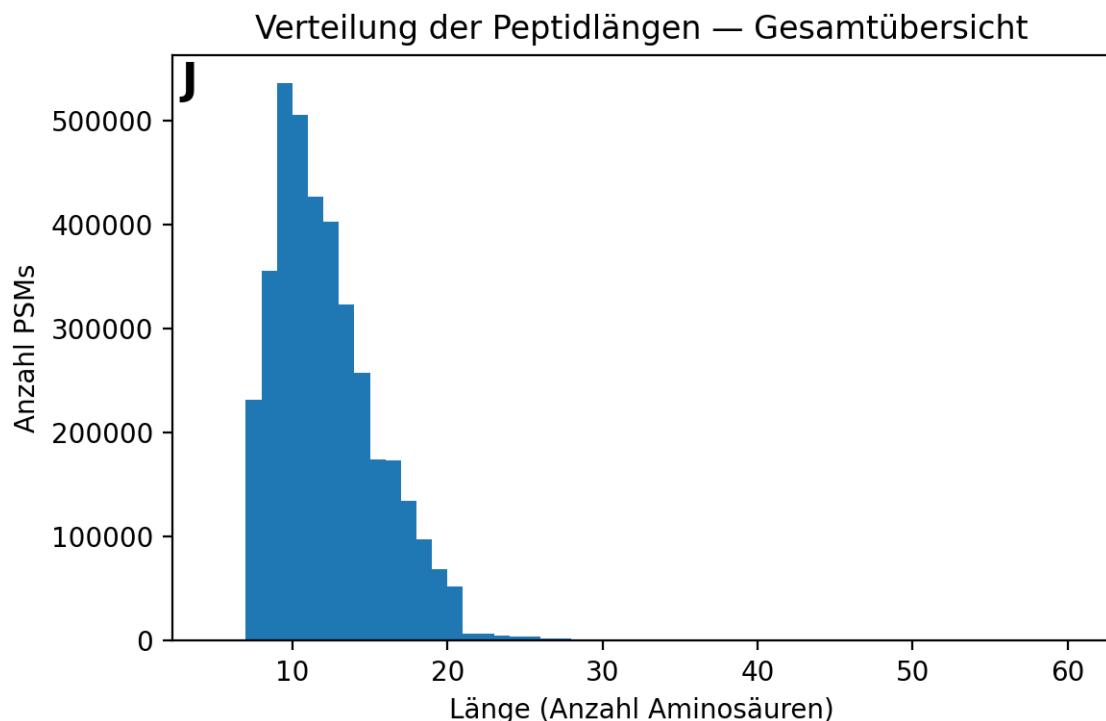


Figure S1: Histogram of peptide lengths for all data

5.2 Zusätzliche Tabellen

Weitere Tabellen...

Dies ist ein Mock-Beispiel.

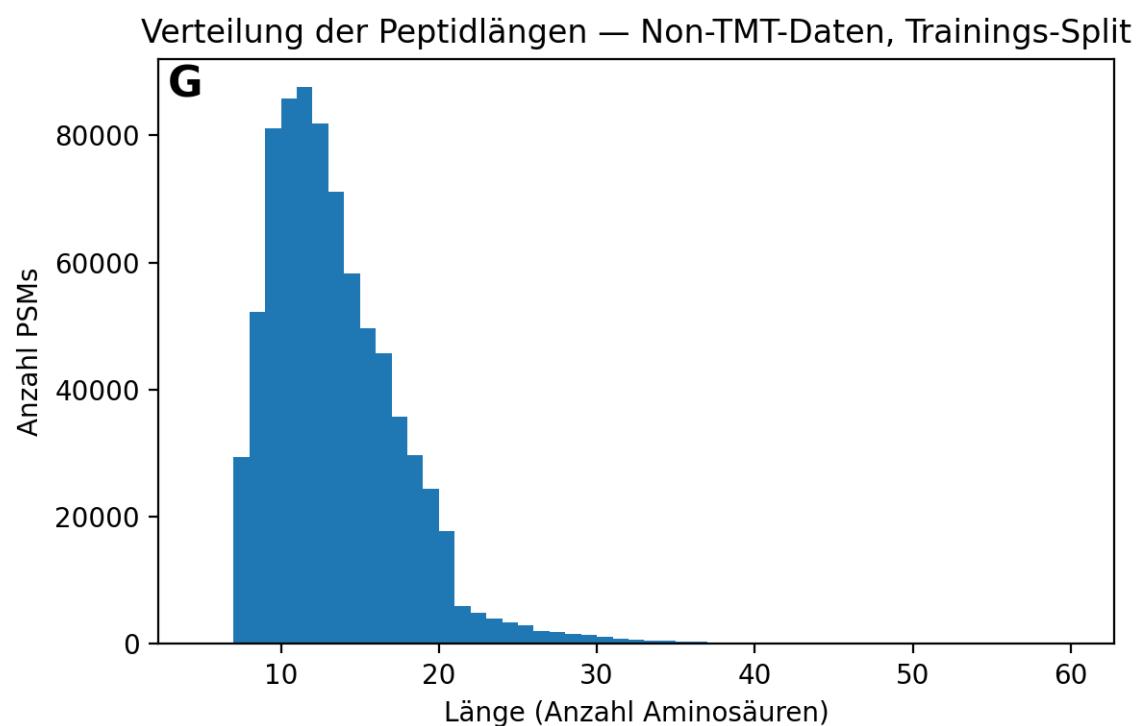


Figure S2: Histogram of peptide lengths for training data (TMT false)

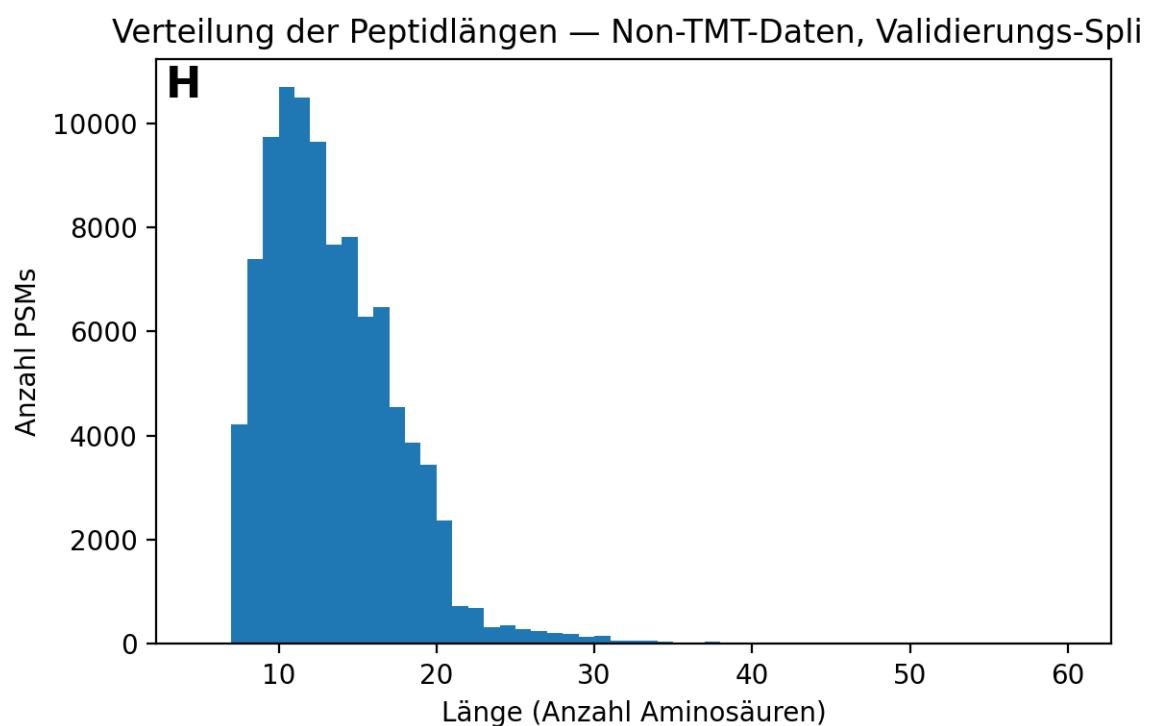


Figure S3: Histogram of peptide lengths for validation data (TMT false)